

## Simple Regression

```
In [3]: import pandas as pd
import numpy as np
import os
from sklearn.model_selection import train_test_split, cross_val_score, cross_val_predict
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [4]: df = pd.read_csv("advertising.csv", usecols = [1,2,3,4])
```

```
In [8]: df.head()
```

Out[8]:

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

```
In [9]: df = df.iloc[:,1:len(df)]
```

```
In [113]:
```

Out[113]:

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4

	TV	radio	newspaper	sales
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

```
In [10]: df.describe().T
```

```
Out[10]:
```

	count	mean	std	min	25%	50%	75%	max
radio	200.0	23.2640	14.846809	0.0	9.975	22.90	36.525	49.6
newspaper	200.0	30.5540	21.778621	0.3	12.750	25.75	45.100	114.0
sales	200.0	14.0225	5.217457	1.6	10.375	12.90	17.400	27.0

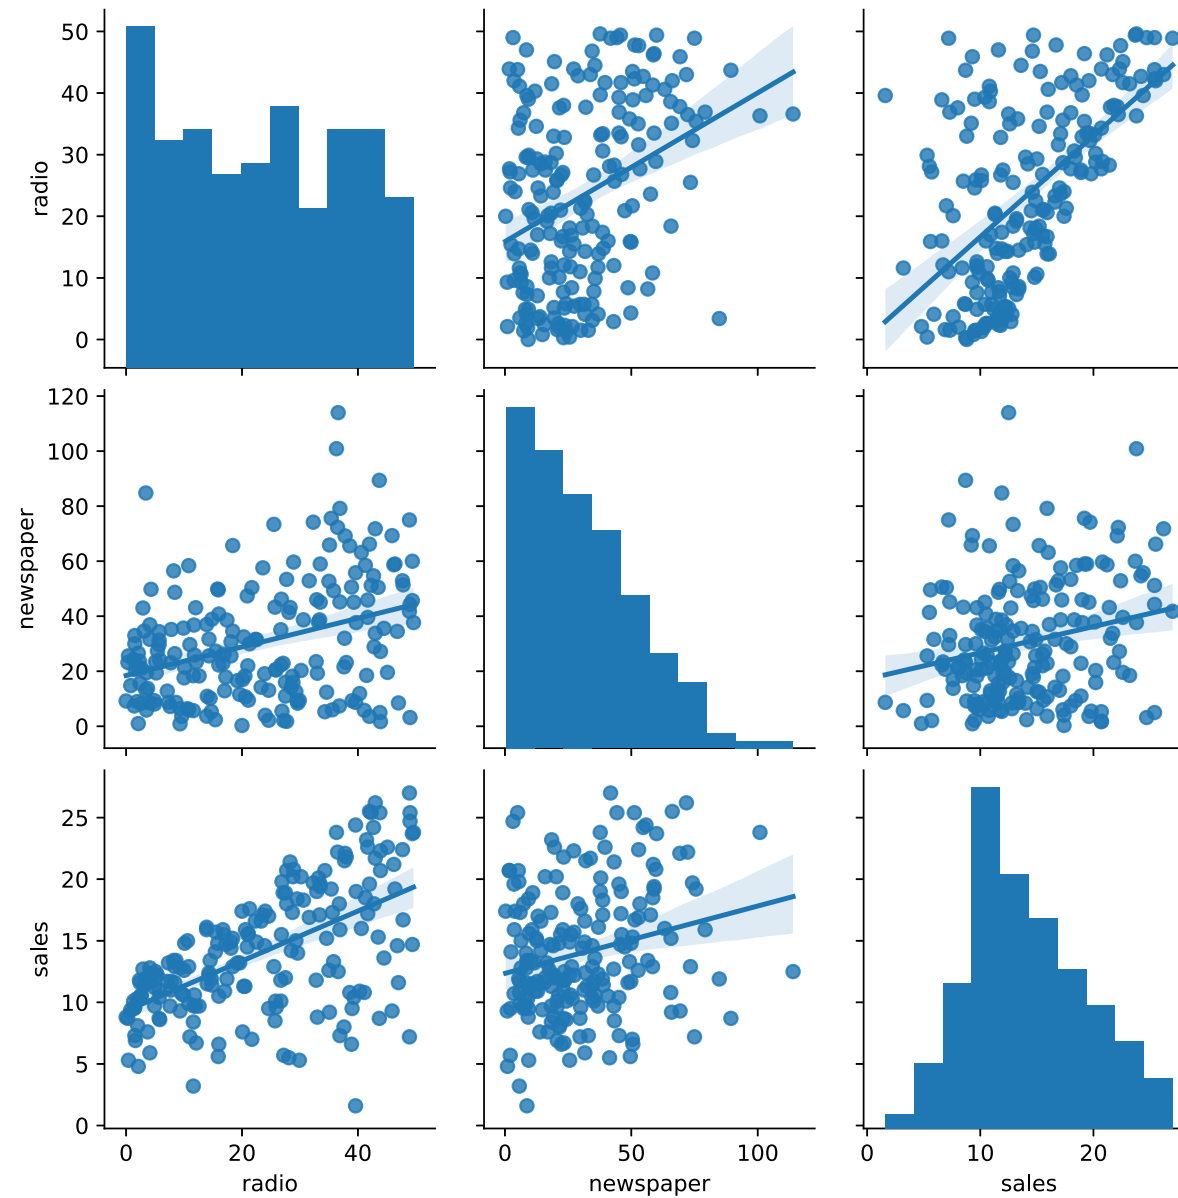
```
In [11]: df.corr()
```

```
Out[11]:
```

	radio	newspaper	sales
radio	1.000000	0.354104	0.576223
newspaper	0.354104	1.000000	0.228299
sales	0.576223	0.228299	1.000000

```
In [12]: import seaborn as sns
sns.pairplot(df, kind = "reg")
```

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x2f0e04a9390>
```



## Modelling with Stats

```
In [36]: X = df[["TV"]]
```

```
In [15]: import statsmodels.api as sm
```

```
In [37]: X = sm.add_constant(X)
```

```
In [38]: y = df["sales"]
```

```
In [39]: lm = sm.OLS(y,X)
```

```
In [40]: model = lm.fit()
```

```
In [41]: model.summary() #coef and TV kısmında tv nin 1 biriminin satış üzerinde ki etkisi
```

Out[41]: OLS Regression Results

<b>Dep. Variable:</b>	sales		<b>R-squared:</b>	0.612		
<b>Model:</b>	OLS		<b>Adj. R-squared:</b>	0.610		
<b>Method:</b>	Least Squares		<b>F-statistic:</b>	312.1		
<b>Date:</b>	Sat, 27 Jun 2020		<b>Prob (F-statistic):</b>	1.47e-42		
<b>Time:</b>	14:35:20		<b>Log-Likelihood:</b>	-519.05		
<b>No. Observations:</b>	200		<b>AIC:</b>	1042.		
<b>Df Residuals:</b>	198		<b>BIC:</b>	1049.		
<b>Df Model:</b>	1					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>const</b>	7.0326	0.458	15.360	0.000	6.130	7.935
<b>TV</b>	0.0475	0.003	17.668	0.000	0.042	0.053

<b>Omnibus:</b>	0.531	<b>Durbin-Watson:</b>	1.935
<b>Prob(Omnibus):</b>	0.767	<b>Jarque-Bera (JB):</b>	0.669
<b>Skew:</b>	-0.089	<b>Prob(JB):</b>	0.716
<b>Kurtosis:</b>	2.779	<b>Cond. No.</b>	338.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Another method

```
In [5]: import statsmodels.formula.api as smf
lm = smf.ols("sales ~ TV", df)
model = lm.fit()
model.summary()
```

Out[5]: OLS Regression Results

Dep. Variable:	sales	R-squared:	0.612			
Model:	OLS	Adj. R-squared:	0.610			
Method:	Least Squares	F-statistic:	312.1			
Date:	Sat, 04 Jul 2020	Prob (F-statistic):	1.47e-42			
Time:	12:45:26	Log-Likelihood:	-519.05			
No. Observations:	200	AIC:	1042.			
Df Residuals:	198	BIC:	1049.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	7.0326	0.458	15.360	0.000	6.130	7.935

<b>TV</b>	0.0475	0.003	17.668	0.000	0.042	0.053
<b>Omnibus:</b>	0.531	<b>Durbin-Watson:</b>	1.935			
<b>Prob(Omnibus):</b>	0.767	<b>Jarque-Bera (JB):</b>	0.669			
<b>Skew:</b>	-0.089	<b>Prob(JB):</b>	0.716			
<b>Kurtosis:</b>	2.779	<b>Cond. No.</b>	338.			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [43]: `model.conf_int()`

Out[43]:

	0	1
<b>Intercept</b>	6.129719	7.935468
<b>TV</b>	0.042231	0.052843

In [44]: `model.f_pvalue`

Out[44]: 1.4673897001947095e-42

In [45]: `print("f_pvalue:", "%.4f"% model.fvalue)`

f\_pvalue: 312.1450

In [46]: `print("t_pvalue:", "%.4f"% model.tvalues[0:1])`

t\_pvalue: 15.3603

In [47]: `model.mse_model`

Out[47]: 3314.6181668686486

```
In [48]: model.rsquared_adj
```

```
Out[48]: 0.6099148238341623
```

```
In [49]: model.fittedvalues[0:5]
]
```

```
Out[49]: 0    17.970775
         1     9.147974
         2     7.850224
         3    14.234395
         4    15.627218
         dtype: float64
```

```
In [29]: import seaborn as sns
         y[0:5]
```

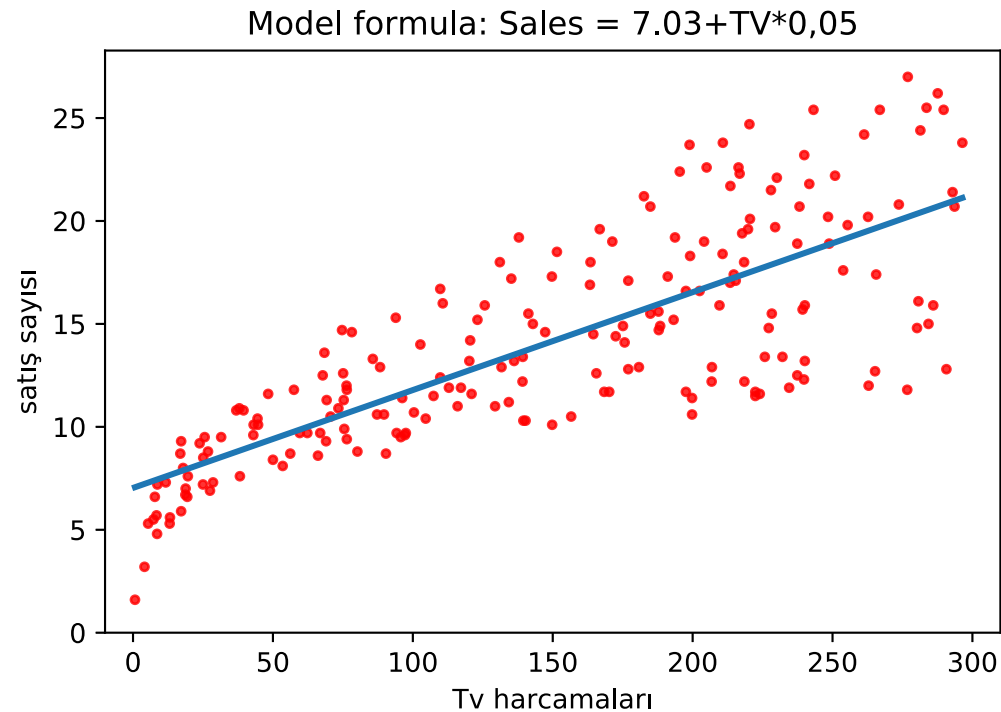
```
Out[29]: 0    22.1
         1    10.4
         2     9.3
         3    18.5
         4    12.9
         Name: sales, dtype: float64
```

```
In [50]: print("sales = " + str("%.2f"% model.params[0])+"TV"+"*" + str("%.2f"%model.params[1]))
```

```
sales = 7.03+TV*0.05
```

```
In [16]: g = sns.regplot(df["TV"], df["sales"], ci =None, scatter_kws={"color":
"r", "s":9})
g.set_title("Model formula: Sales = 7.03+TV*0,05")
g.set_ylabel("satış sayısı")
g.set_xlabel("Tv harcamaları")
plt.xlim(-10,310)
plt.ylim(bottom = 0)
```

```
Out[16]: (0, 28.278347130242828)
```



```
In [9]: import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [8]: from sklearn.linear_model import LinearRegression
```

```
In [10]: X=df[["TV"]]
```

```
In [11]: y=df["sales"]  
reg = LinearRegression()  
model = reg.fit(X,y)  
model.intercept_  
model.coef_
```

```
Out[11]: array([0.04753664])
```



```
In [61]: model.score(X,y)
```

```
Out[61]: 0.611875050850071
```

Prediction

```
In [12]: X=df[["TV"]]
y=df["sales"]
reg = LinearRegression()
model = reg.fit(X,y)
```

```
In [75]: model.predict([[30]])
```

```
Out[75]: array([8.45869276])
```

```
In [13]: k_t = pd.DataFrame({"real_y": y[0:10],
                             "predicted_y": reg.predict(X)[0:10]})
```

```
In [66]: k_t["hata"] = k_t["gercek_y"]-k_t["tahmin_y"]
```

```
In [14]: k_t
```

```
Out[14]:
```

	real_y	predicted_y
0	22.1	17.970775
1	10.4	9.147974
2	9.3	7.850224
3	18.5	14.234395
4	12.9	15.627218
5	7.2	7.446162
6	11.8	9.765950

	real_y	predicted_y
7	13.2	12.746498
8	4.8	7.441409
9	10.6	16.530414

```
In [68]: k_t["hata_kare"] = k_t["hata"]**2
```

```
In [69]: k_t
```

```
Out[69]:
```

	gercek_y	tahmin_y	hata	hata_kare
0	22.1	17.970775	4.129225	17.050503
1	10.4	9.147974	1.252026	1.567569
2	9.3	7.850224	1.449776	2.101851
3	18.5	14.234395	4.265605	18.195390
4	12.9	15.627218	-2.727218	7.437719
5	7.2	7.446162	-0.246162	0.060596
6	11.8	9.765950	2.034050	4.137358
7	13.2	12.746498	0.453502	0.205664
8	4.8	7.441409	-2.641409	6.977040
9	10.6	16.530414	-5.930414	35.169814

```
In [70]: np.sum(k_t["hata_kare"])
```

```
Out[70]: 92.90350329638103
```

```
In [82]: X = df.drop("sales", axis = 1)
y = df["sales"]
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.30,
```

```
train_size =  
0.70, random_state =0)
```

```
In [83]: training = df.copy()
```

```
In [84]: lm = sm.OLS(y_train,X_train)
```

```
In [86]: model = lm.fit()  
model.summary()
```

Out[86]:

OLS Regression Results

<b>Dep. Variable:</b>	sales	<b>R-squared (uncentered):</b>	0.984			
<b>Model:</b>	OLS	<b>Adj. R-squared (uncentered):</b>	0.984			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2862.			
<b>Date:</b>	Sat, 27 Jun 2020	<b>Prob (F-statistic):</b>	2.50e-123			
<b>Time:</b>	16:40:24	<b>Log-Likelihood:</b>	-289.37			
<b>No. Observations:</b>	140	<b>AIC:</b>	584.7			
<b>Df Residuals:</b>	137	<b>BIC:</b>	593.6			
<b>Df Model:</b>	3					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>TV</b>	0.0520	0.002	33.691	0.000	0.049	0.055
<b>radio</b>	0.2321	0.010	22.911	0.000	0.212	0.252
<b>newspaper</b>	0.0188	0.008	2.371	0.019	0.003	0.034
<b>Omnibus:</b>	1.996	<b>Durbin-Watson:</b>	1.890			
<b>Prob(Omnibus):</b>	0.369	<b>Jarque-Bera (JB):</b>	1.965			
<b>Skew:</b>	0.284	<b>Prob(JB):</b>	0.374			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]: ### Scikit-learn model
```

```
In [87]: lm = LinearRegression()  
model = lm.fit(X_train, y_train)
```

```
In [88]: model.intercept_
```

```
Out[88]: 2.880255286331323
```

```
In [89]: model.coef_
```

```
Out[89]: array([0.04391531, 0.20027962, 0.00184368])
```

```
In [93]: rmse = np.sqrt(mean_squared_error(y_train,model.predict(X_train)))## ha  
ta oranı
```

```
In [94]: rmse
```

```
Out[94]: 1.5768437866753109
```

```
In [97]: cross_val_score(model,X,y, cv = 10, scoring = "r2").mean()
```

```
Out[97]: 0.8853562237979616
```

```
In [101]: np.sqrt(-cross_val_score(model,X_train,y_train, cv = 10, scoring = "neg  
_mean_squared_error")).mean() ##validating and real result
```

```
Out[101]: 1.59979865949965
```

PLC

```
In [136]: hit = pd.read_csv("hitters.csv")
```

```
In [137]: df1 = hit.copy()
```

```
In [138]: df1 = df1.dropna()  
df1.head()
```

Out[138]:

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAAtBat	CHits	CHmRun	CRuns	CRBI	CWal
1	315	81	7	24	38	39	14	3449	835	69	321	414	3
2	479	130	18	66	72	76	3	1624	457	63	224	266	2
3	496	141	20	65	78	37	11	5628	1575	225	828	838	3
4	321	87	10	39	42	30	2	396	101	12	48	46	
5	594	169	4	74	51	35	11	4408	1133	19	501	336	1

```
In [139]: dms = pd.get_dummies(df1[["League", "Division", "NewLeague"]])  
dms.head()
```

Out[139]:

	League_A	League_N	Division_E	Division_W	NewLeague_A	NewLeague_N
1	0	1	0	1	0	1
2	1	0	0	1	1	0
3	0	1	1	0	0	1
4	0	1	1	0	0	1
5	1	0	0	1	1	0

```
In [140]: y = df1["Salary"]
```

```
In [141]: X_ = df1.drop(["Salary", "League", "Division", "NewLeague"], axis = 1).astype("float64") #we removed dependent and categoric variables from dataset
```

```
In [142]: X = pd.concat([X_, dms[["League_N", "Division_W", "NewLeague_N"]]], axis = 1)
```

```
In [143]: X
```

Out[143]:

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAAtBat	CHits	CHmRun	CRuns	CRBI
1	315.0	81.0	7.0	24.0	38.0	39.0	14.0	3449.0	835.0	69.0	321.0	414.0
2	479.0	130.0	18.0	66.0	72.0	76.0	3.0	1624.0	457.0	63.0	224.0	266.0
3	496.0	141.0	20.0	65.0	78.0	37.0	11.0	5628.0	1575.0	225.0	828.0	838.0
4	321.0	87.0	10.0	39.0	42.0	30.0	2.0	396.0	101.0	12.0	48.0	46.0
5	594.0	169.0	4.0	74.0	51.0	35.0	11.0	4408.0	1133.0	19.0	501.0	336.0
...	...	...	...	...	...	...	...	...	...	...	...	...
317	497.0	127.0	7.0	65.0	48.0	37.0	5.0	2703.0	806.0	32.0	379.0	311.0
318	492.0	136.0	5.0	76.0	50.0	94.0	12.0	5511.0	1511.0	39.0	897.0	451.0
319	475.0	126.0	3.0	61.0	43.0	52.0	6.0	1700.0	433.0	7.0	217.0	93.0
320	573.0	144.0	9.0	85.0	60.0	78.0	8.0	3198.0	857.0	97.0	470.0	420.0
321	631.0	170.0	9.0	77.0	44.0	31.0	11.0	4908.0	1457.0	30.0	775.0	357.0

263 rows × 19 columns



```
In [144]: X_train, X_test, y_train, y_test = train_test_split(X,
                                                             y,
                                                             test_size = 0.25,
                                                             random_state = 42)
```

```
In [145]: print("X_train", X_train.shape)
          print("y_train", y_train.shape)
```

```
print("X_test",X_test.shape)
print("y_test",y_test.shape)

training = df.copy()

print("training", training.shape)
```

```
X_train (197, 19)
y_train (197,)
X_test (66, 19)
y_test (66,)
training (200, 4)
```

```
In [146]: from sklearn.decomposition import PCA
          from sklearn.preprocessing import scale
          pca= PCA()
```

```
In [147]: X_reducen_train = pca.fit_transform(scale(X_train))
```

```
In [148]: np.cumsum(np.round(pca.explained_variance_ratio_,decimals = 4)*100)[0:
          100]
```

```
Out[148]: array([38.18, 59.88, 70.88, 78.88, 84.18, 88.45, 92.05, 94.86, 96.34,
                 97.28, 98.01, 98.68, 99.18, 99.49, 99.74, 99.9 , 99.96, 99.98,
                 99.99])
```

```
In [149]: lm = LinearRegression()
```

```
In [150]: pcr_model = lm.fit(X_reducen_train,y_train)
```

```
In [151]: pcr_model.intercept_
```

```
Out[151]: 543.4834416243655
```

```
In [154]: y_pred = pcr_model.predict(X_reducen_train)
```

```
In [156]: y_pred[0:5]
```

```
Out[156]: array([377.44484744, 802.19452124, 495.60987745, 112.53177731,  
                426.21613066])
```

```
In [158]: np.sqrt(mean_squared_error(y_train, y_pred))
```

```
Out[158]: 289.3292825564976
```

```
In [15]: ! jupyter nbconvert --to html MachineLearning1.ipynb
```

```
[NbConvertApp] Converting notebook MachineLearning1.ipynb to html  
[NbConvertApp] Writing 685549 bytes to MachineLearning1.html
```

```
In [ ]:
```

```
In [ ]:
```